

Blue algae project report

Nadezhda Novikovskaia, Ossi Räisä, Janani Subramanian

October 22, 2019

1 Data

1.1 Healthcare visits

The healthcare visits data was downloaded from the THL API. The API return logged visits to healthcare providers for each week of 2018 and 2019 (until week 38). To find to locations of the providers, a list of subproviders with their municipalities was used. We assumed that each provider is in the municipality where most of their subproviders are. Then the providers were grouped by their municipality and the groups were summed with missing values assumed to be zeros giving healthcare visits for each municipality for each week.

Figures 1 and 2 show the visits and visits per population for all municipalities included in the analysis. The drop during the last week may be due to the data being collected in the middle of the week. We were not able to explain the drop when the year changes.




Figure 1: Healthcare visits for all municipalities.




Figure 2: Healthcare visits per population in municipalities.

1.2 Algae

The algae data is from an API used by the Lake and Sea wiki algae table page. The URL for the API was found by looking at the requests made by the page and finding the request returning the algae data.

The API returns the data as a json string containing a list of measurements. From each measurement the date, place and algae level were extracted. The measurements were first compiled to a table with weeks as columns, measurement stations as rows and algae levels as values. There was mostly one measurement per week for each station, so the few conflicting measurements for were ignored. Then the stations were grouped by their municipality and the algae levels of each group were averaged, giving a table with municipalities as rows.

The algae level is a number from 0 to 3 where 0 is no algae and 3 is the most algae. The time period is weeks from 23 to 39 in 2018 and 23 to 37 in 2019. The data has a lot of missing values, even for municipalities.

1.3 Weather and airquality

The weather and airquality data was downloaded from the WFS API of the Finnish Meteorological Institute. The airquality is from the stored query “urban::observations::airquality::hourly::timevaluepair” and the rain and temperature is from “fmi::observations::weather::daily::timevaluepair”. As the airquality returns hourly data getting all of it for two years takes over a hundred queries, it and the weather data was only retrieved for municipalities with no missing algae values and a population over 10 000 people. The airquality data includes several measures for different substances that pollute the air and an index, which was the only value we used.

The airquality, temperature and rain values for were grouped by the week

and year they are from and averaged. There were 22 municipalities that had the algae, weather and airquality data and were large enough.

1.4 Population

The population data was downloaded from Statistics Finland interface services Municipal key figures (2018 population) and Preliminary population structure by area (2019 population).

2 Analysis

To predict healthcare visits we made the data into a single table where each row is a single week of one municipality. The columns are the algae, temperature, rain and airquality values for that week and several previous weeks, the number of which was varied. The table also had the population of the municipality for the row's year and finally the number of healthcare visits for the week and municipality. For the weeks outside the period where algae data is gathered we assumed that no algae was present.

At first we ran linear regression, random forest regression and support vector machine regression on the data for just one municipality without population. The results were not good, so we included more municipalities and their population. This had better results. The numbers shown on the pitch were with 6 municipalities. After the pitch we downloaded the weather and airquality data for the rest of the municipalities and ran the regression methods. As the SVM was not performing well it was replaced by a decision tree.

The first experiment has a total of 4 weeks for each of the features varying with time and the population, for a total of 17 features. With 22 cities the dataset has 1980 rows. The random forest has 100 trees with 10 maximum depth and the decision tree has 3 maximum depth to allow inspecting it.

Table 1 shows the results. Surprisingly even the very simple decision tree gets a high score. Figure 3 shows the actual and predicted visits for Vaasa. The prediction almost constant and does not capture the changes in visits.

Model	Training R^2	Test R^2
Linear Regression	0.783	0.794
Random Forest	0.955	0.950
Decision Tree	0.911	0.922

Table 1: Results of the first experiment




Figure 3: Healthcare visits in Vaasa. The prediction is very constant and does not capture the changes in visits. The change of prediction in 2019 is a result of the population changing.

Figure 4 shows the decision tree. As we can see $X[16]$, that being the population, is used in almost all of the nodes. The linear regression coefficients are -701, -919, 151 and -340 for the algae on weeks 0, -1, -2, and -3 respectively, 0.24 for population and 0 for the rest. The algae coefficients are larger than the population because the algae level is from 0 to 3 while the population is over 10 000.

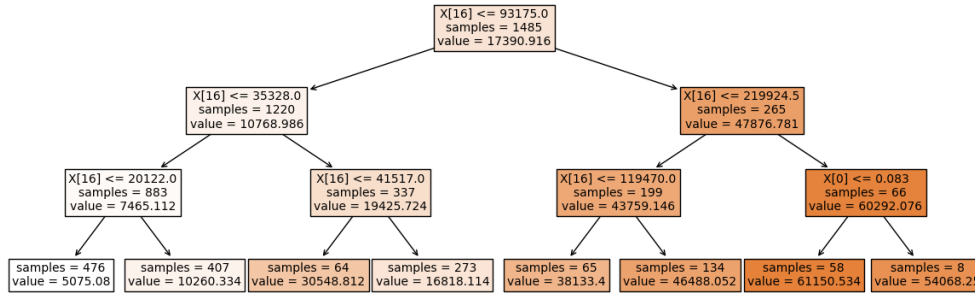


Figure 4: $X[16]$ is the population and $X[0]$ is the algae level for the week in question. Most nodes only use the population.

This means that the population actually explains most of the healthcare visits, which is to be expected. To see if the other features are useful at all, for the next experiment we removed all features except population. Model parameters were not changed.

Table 2 shows the results. The test scores show very minor drops for random

forest and decision tree and no change for linear regression, meaning that the other features are not very useful in the prediction.

Model	Training R^2	Test R^2
Linear Regression	0.783	0.794
Random Forest	0.944	0.949
Decision Tree	0.910	0.920

Table 2: Results of the second experiment

For the final experiment we replace healthcare visits with visits per population and removed population from the features. Table 3 shows the results.

Model	Training R^2	Test R^2
Linear Regression	0.017	-0.003
Random Forest	0.097	-0.021
Decision Tree	0.026	-0.001

Table 3: Results of the third experiment

The scores are very bad so the features are not able to predict the healthcare visits per population. As the training score is low even for the random forest we tried increasing the maximum depth of the it which, but it did not change the scores significantly. Figure 5 shows the actual and predicted visits per population in Vaasa. Again, the prediction is constant and this time it is also completely off.

3 Conclusion

Our experiments show that population is a good predictor of healthcare visits, which is to be expected. The other included features are not good predictors. This is likely due to people going to the doctors for many different reasons, with blue algae accounting for a small percentage of them. Limiting the visits to only include symptoms or diagnoses related to blue algae could allow to find a correlation, but unfortunately THL does not have open data on a weekly timespace that could be used.




Figure 5: Healthcare visits per population in Vaasa. The prediction is still constant and does not capture the changes in visits. It is also in a completely wrong place.